

Gianluca Sartori



Monitoring SQL Server Without Breaking the Bank

Organizers



Sponsors



Gianluca Sartori

Independent SQL Server consultant

SQL Server MVP, MCTS, MCITP, MCT

Works with SQL Server since version 7

DBA @ Scuderia Ferrari

Blog: spaghettidba.com

Twitter: [@spaghettidba](https://twitter.com/spaghettidba)



Agenda

- What we need
 - What the market offers
 - What SQL Server offers
 - What the community offers
- The TICK stack
- How you do it

A large, teal-colored abstract graphic on the left side of the slide, consisting of several thick, curved lines that form a stylized, open shape.

What we need



Monitoring

- Troubleshooting

What happened tonight at 03:40?

- Tune performance

Which queries consume most resources?

- Capacity Planning

How fast are my databases growing?

- Baselineing

Is the system behaving normally?

- Alerting

Hey, look: something's wrong here!



#SQLSAT895



What the market offers

Buy a product

- Big value
- Big money \$\$

SentryOne

solarwinds 

Quest

I D E R A

 ApexSQL


redgate

What SQL Server offers

Use SQL Server built-in tools

- SQLTrace
- Extended Events
- Alerts
- Activity Monitor
- SSMS Reports
- Data Collector



SQL Server Data Collector

PROS:

- Included in SQL Server
- Offers some reports
- Good for troubleshooting

CONS:

- Unreliable
- Slow
- No alerting
- Hard to customize/extend
- Deprecated

What the community offers

Generic monitoring tools

Complete solutions

- Nagios
- Zabbix
- Cacti

Tools

- Elasticsearch
- Kibana

SQL Server monitoring tools

Complete solutions

- SQLWatch
- TICK stack

Tools

- dbatools
- ExtendedTSQLCollector
- XESmartTarget

A large, teal-colored abstract graphic on the left side of the slide, consisting of several thick, curved lines that form a stylized, open shape.

The TICK stack



The TICK stack



Open source
Cross platform



telegraf



influxdb



chronograf



Kapacitor



#SQLSAT895





Time series database

Written in GO

Designed to store time series data

Offers two query languages

- InfluxQL - similar to SQL
- Flux - inspired by Javascript / LINQ





Data collection agent

Can monitor several information sources:

- Operating system
- Databases
- Network
- Message Queues
- Apps
- ...

Works with *plugins*

Over 100 plugins

- Apache
- Cassandra
- Docker
- Elasticsearch
- Graylog
-
- SQL Server!





Administration interface for Influxdb

User Interface for Kapacitor

Can be used to explore data

Can create dashboards to visualize data

Can set up and manage alerts in Kapacitor





Data processing engine

Can publish data to topics

Handlers subscribe to topics and perform actions

Alerts can be sent to several channels, using *plugins*

- Email
- Slack
- ...





Monitoring / analytics platform

Native Influxdb support

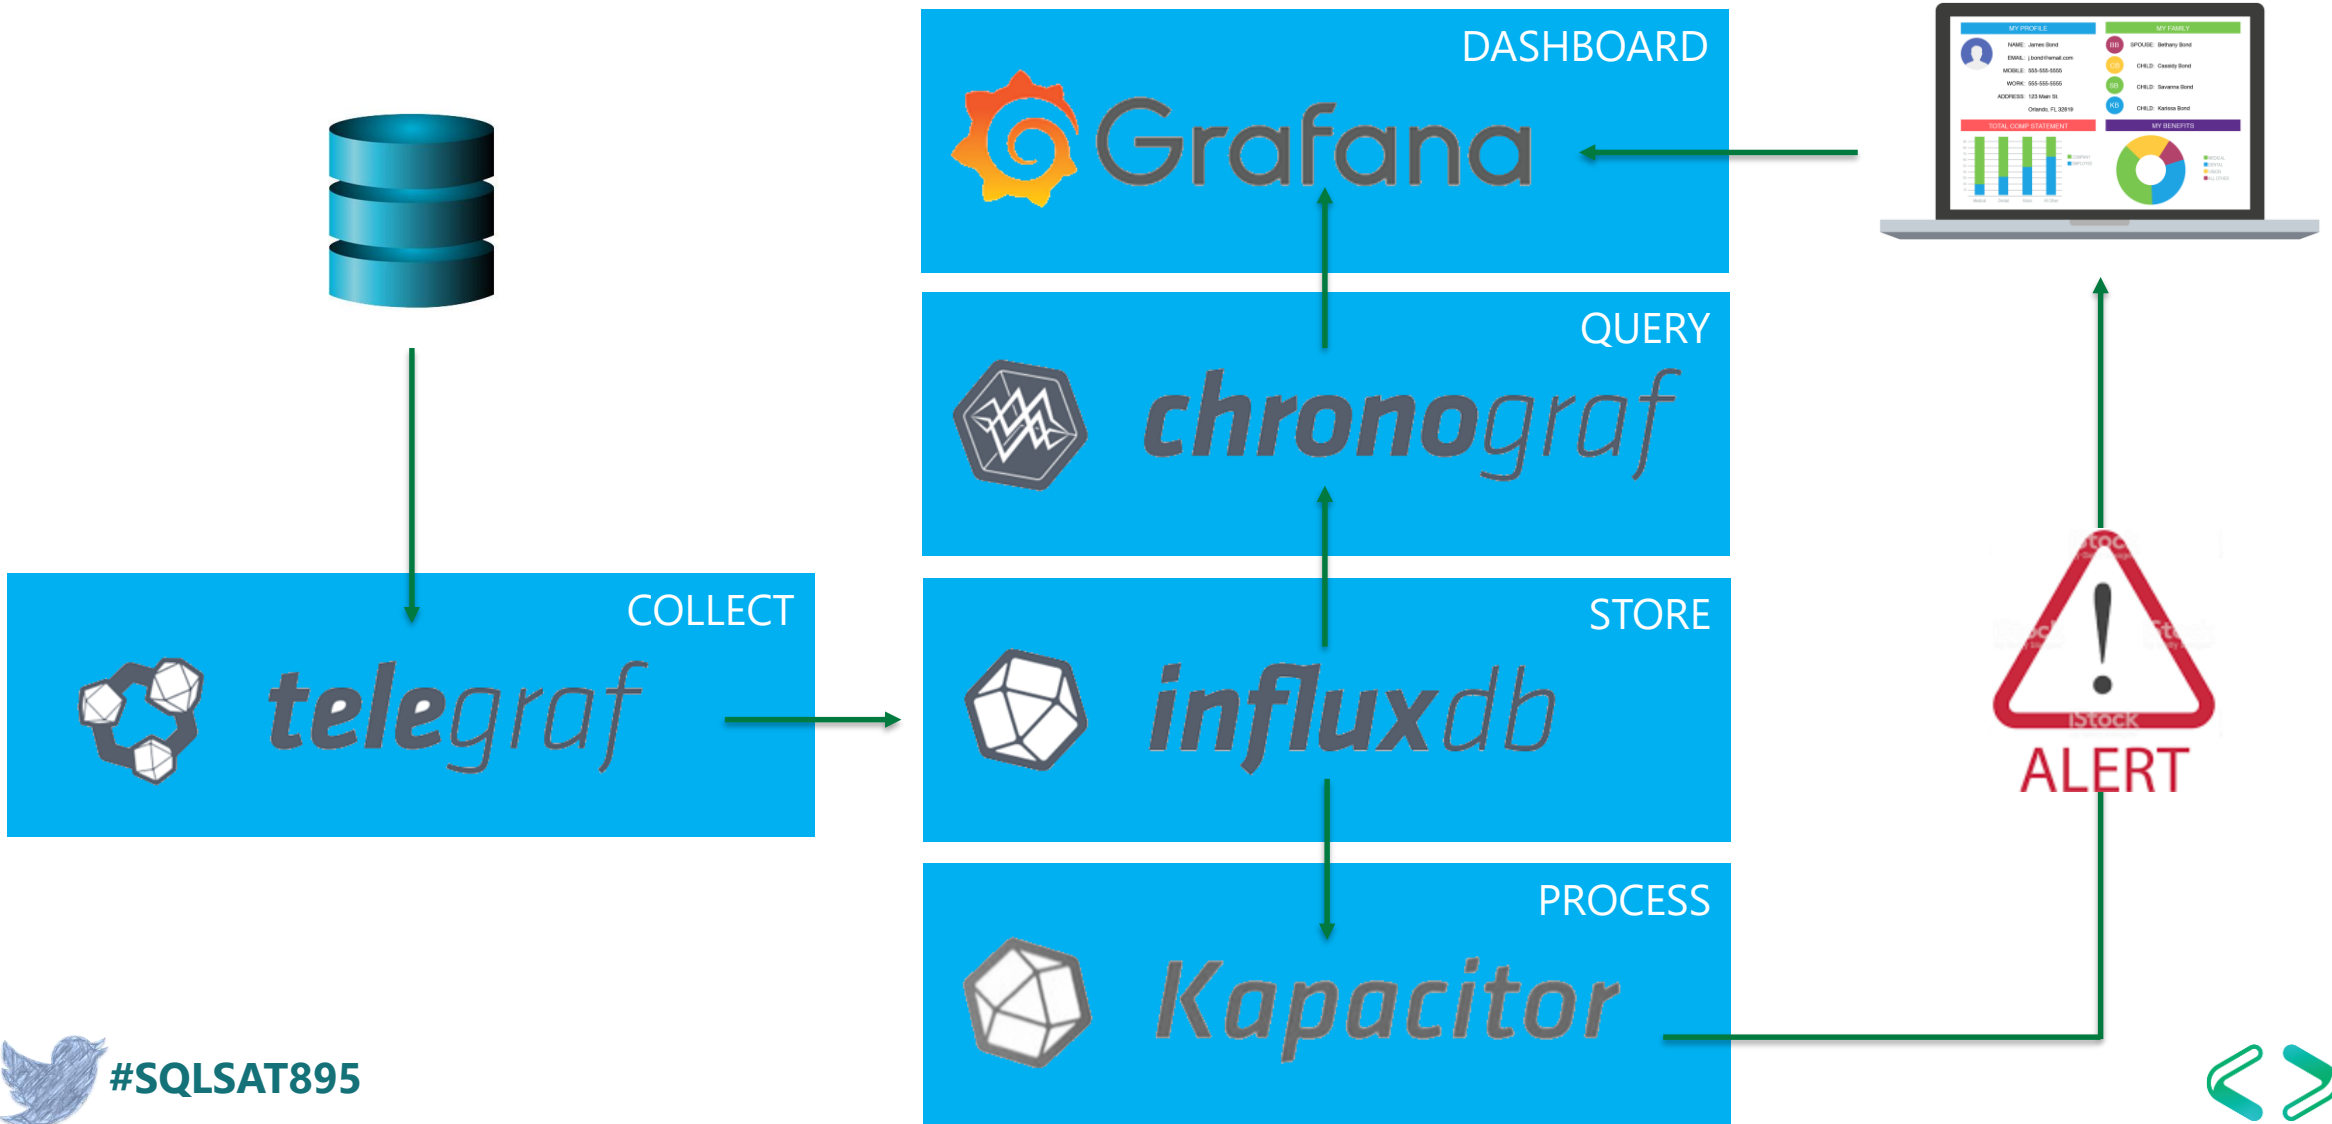
More data sources supported via *plugins*

Easy dashboard creation

Multiple built-in / community dashboards



Architecture



A large, teal-colored abstract graphic on the left side of the slide, consisting of several thick, curved lines that form a stylized, open shape.

How you do it





1. Download Influxdb
2. Unzip the archive (c:\tick maybe?)
3. Run `influxD.exe`
 - Open port in FW (default 8086)
 - Run it as a service maybe?



NSSM - the Non-Sucking Service Manager

Ever used srvany? You get why NSSM then...

1. Download NSSM
2. Unzip the archive (c:\tick maybe?)
3. `nssm install influxdb`





1. Download Telegraf
2. Unzip the archive (`c:\tick` maybe?)
3. Generate and edit the config file
 - Log
 - Input
 - Output



1. Download Chronograf
2. Unzip the archive
3. Install as a service
4. Open in browser
 - Create a connection
 - Create dashboards
 - Explore the data



Influx DB entity

Measurements

Tags

Fields

Points

SQL DB entity

Tables

Columns (indexed)

Columns (not indexed)

Rows

- **Tags** are indexed, and are used **to filter and aggregate data**
=> view CPU usage (the Field) **by host** (the context)
- **Fields** are not indexed and should **only be aggregated**.
A field should **represent a metric**.
=> CPU usage





1. Download Grafana
2. Unzip the archive
3. Install as a service
4. Open in browser
 - Create a connection to Influxdb
 - Create (or download) dashboards
 - Explore the data





1. Configure alerts by dashboard
2. Multiple channels available
 - Email
 - Slack
 - Teams
 - Telegram
 - ...

New Notification Channel

Name	
Type	Email
Default (send on all alerts)	Pushover webhook
Include image	HipChat Kafka REST Proxy
Disable Resolve Message	Sensu Telegram
Send reminders	VictorOps DingDing
Email addresses	Email Google Hangouts Chat PagerDuty Slack Threema Gateway Prometheus Alertmanager Discord LINE OpsGenie Microsoft Teams

You can enter multiple email addresses using a ";" separator



Continuous Queries

- Scheduled queries in InfluxQL
- Provide downsampling and aggregation of data
- Can write to a different retention policy

```
CREATE CONTINUOUS QUERY "cq_basic" ON "transportation"  
BEGIN  
    SELECT mean("passengers") INTO "average_passengers"  
    FROM "bus_data" GROUP BY time(1h)  
END
```



#SQLSAT895



TICK stack: recap

- ✓ Troubleshooting

Grafana dashboards, standard sqlserver dashboard + custom dashboards

- ✓ Tune performance

Grafana dashboards, standard sqlserver dashboard + custom dashboards

- ✓ Capacity Planning

Influxdb continuous queries, custom Grafana dashboards

- ✓ Baselining

Influxdb continuous queries, custom Grafana dashboards

- ✓ Alerting

Kapacitor or Grafana

Resources

- Influxdb, Telegraf, Chronograf, Kapacitor
- <https://www.influxdata.com/>
- NSSM - Non Sucking Service Manager
- <https://nssm.cc>
- Grafana
- <https://grafana.com>



Thanks!



#SQLSAT895